

Enterprise AI Agent Platform

Product Vision & North Star

Document Status: Product North Star

Audience: Engineering, AI/ML

Purpose: Define exactly what we are building before implementation begins.

1. The Vision

We are building an **Enterprise AI Agent Platform** that allows organizations to create, customize, deploy, govern and operate AI agents without needing to build every AI workflow from scratch.


The platform should become the place where an enterprise can go from:

“I have a business problem.”

to:

“I have a production-ready AI agent solving this problem.”

The platform should support the complete lifecycle:



Test & Evaluate
↓
Approve
↓
Deploy
↓
Run in Production
↓
Monitor
↓
Improve

This is not just an AI chatbot.

This is not just a workflow builder.

This is not just an LLM wrapper.

This is an **enterprise agent operating platform**.

2. The Problem We Are Solving

Today, companies want to use AI for business operations.

For example:

- Finance wants an invoice agent.
- HR wants an employee onboarding agent.
- Sales wants a lead qualification agent.
- Customer support wants a resolution agent.
- IT wants an incident management agent.
- Legal wants a contract analysis agent.
- Operations wants an approval automation agent.

But creating these systems currently requires:

LLM expertise
+
Backend engineering
+
API integrations
+
RAG
+
Workflow orchestration
+
Security
+
Authentication
+
Monitoring
+
Evaluation
+
Deployment

This makes AI automation expensive and slow.

Our platform removes much of that complexity.

3. The Core Product Promise

A customer should be able to say:

“I want an agent that checks our invoices every morning, identifies overdue invoices, analyzes them, and notifies the finance team.”

And the platform should transform that request into a working agent.

The experience should be:

USER

"Create an invoice collection agent."



AI UNDERSTANDS REQUIREMENT



AI DESIGNS AGENT



VISUAL WORKFLOW APPEARS



USER REVIEWS / MODIFIES



CONNECTS TO SYSTEMS



TESTS AGENT



DEPLOYS



AGENT RUNS IN PRODUCTION

The user should not need to understand:


- LangGraph
- MCP

- vector databases
- embeddings
- tool calling
- queues
- workflow engines
- model APIs

Those are implementation details of our platform.

4. What We Are Actually Building

The product has five major layers.



5. Pillar 1 — Create

The first pillar is the **Agent Builder**.

Customers should have multiple ways to create an agent.

A. Create with AI


User says:

“Create an agent that handles employee onboarding.”

The platform creates the initial workflow.

B. Visual Builder

Users can manually build:



C. Start from Template

User chooses:

Finance Agent
HR Agent
Sales Agent
Support Agent
IT Agent

and customizes it.

6. Pillar 2 — Agent Intelligence

The platform provides the intelligence layer.

An agent can:

- Understand user intent
- Reason about tasks
- Decide which tools to use
- Retrieve enterprise knowledge
- Remember relevant information
- Generate structured outputs
- Ask humans for help
- Handle failures
- Continue long-running tasks

The platform should support both:

Single-agent workflows

User


↓

Agent

↓
Tools

and eventually:

Multi-agent systems



7. Pillar 3 — Connect


An agent becomes useful when it can interact with enterprise systems.

The platform must allow agents to connect to:

CRM
ERP
HRIS
Accounting
Databases
Email
Slack
Microsoft Teams
WhatsApp
Ticketing Systems
Internal APIs
MCP Servers

Example:

Invoice Agent



The customer should be able to connect systems without writing custom integration code for every agent.

8. Pillar 4 — Knowledge

Agents must understand company-specific information.

Example:


An HR agent needs:

HR Policies
Employee Handbook
Country Regulations
Benefits Documents
Leave Policies
Onboarding Documents

The platform therefore provides reusable enterprise knowledge bases.

Knowledge Base





One knowledge base can be used by multiple agents.

9. Pillar 5 — Operate

Creating an agent is only half the product.

The difficult part is operating agents reliably in production.

We need:

- Deploy
- Monitor
- Debug
- Evaluate
- Version
- Rollback
- Audit
- Secure
- Scale

Every production execution should be observable.

Example:

Employee Onboarding Agent

Execution #1234

- ✓ Employee received
- ✓ Documents validated
- ✓ HR policy retrieved
- ✓ Employee record created
- ✓ Approval requested
- ✓ HR approved
- ✓ Welcome email sent


Duration: 8.2 sec


Cost: \$0.09

Status: SUCCESS

10. The Agent Lifecycle

Every agent follows:





This lifecycle is one of the core concepts of the platform.

11. The Most Important User Experience

The platform should feel like:

“Describe → See → Edit → Test → Deploy.”

Not:

Configure 25 technical settings → write code → configure APIs → deploy.

12. Example: Creating an Agent

User opens:

Create Agent

and writes:

Create an employee onboarding agent.

The platform responds:

I designed the following agent:

Employee Onboarding Agent

Trigger



Get Employee Information



Validate Documents



Search HR Policies



Create Employee Record



Human Approval



Send Welcome Message

The user can then say:

Add a manager approval before creating the employee record.

The platform updates the workflow:

Validate Documents



Search HR Policies



Manager Approval



Create Employee Record


The user can visually inspect and modify everything.

13. The Agent Builder Should Feel Alive

The canvas should not be a static diagram.

When the agent runs, users should see activity.

Example:



The user should be able to watch the agent execute.

This is an important part of the product experience.

14. Conversational Editing

The platform should support continuous interaction.

User:

Add a fraud check.

Platform:

Fraud Detection



?

User:

If fraud risk is high, send it to a human.

Platform:

Fraud Detection



Risk > Threshold?

/

YES



Human
Review

\

NO



Continue

User:

Deploy it.

Platform:

Validation ✓

Tests ✓

Security ✓

Ready for deployment.

[Deploy to Production]

The agent builder becomes an **AI development environment**, not simply a drag-and-drop tool.

15. Human-in-the-Loop

We do not want uncontrolled AI automation.

Enterprise agents must know when humans are required.

Example:

```
AI
↓
Analyze loan
↓
Risk score
↓
Human approval
↓
Execute
```

Human approval should support:

- Approver
- Role
- Timeout
- Escalation
- Approval conditions
- Rejection path
- Audit trail

16. Trust Is a Core Product Feature

Enterprise customers will ask:

“What exactly did the agent do?”

We must be able to answer.

For every execution:

What did the user ask?
What did the agent decide?
What model was used?
What knowledge was retrieved?
What tools were called?
What data was sent?
What actions happened?
Who approved it?
What was the result?
How much did it cost?

Therefore, every execution needs a complete trace.

17. Enterprise Governance

Customers must control what agents can do.

Example:

Agent

- |
- |— Can read invoices
- |— Can create invoice
- |— Cannot delete invoice
- |— Can send Teams message
- |— Requires approval for payment

Permissions must exist at:

User
↓
Workspace
↓
Agent

↓
Knowledge
↓
Tool
↓
Action

18. Marketplace

The long-term platform should become a marketplace of business agents.

Examples:

Finance

Invoice Collection Agent
Expense Audit Agent
Financial Reporting Agent
Accounts Receivable Agent
HR

Employee Onboarding Agent
Recruitment Agent
Leave Management Agent
Employee Support Agent
IT

IT Helpdesk Agent
Incident Response Agent
Access Request Agent

Customers can:

Discover
↓
Install

↓
Configure
↓
Test
↓
Deploy

19. Marketplace Philosophy

Marketplace agents are not simple chatbot prompts.

A marketplace agent should include:

Agent
+
Workflow
+
Tools
+
Knowledge
+
Guardrails
+
Evaluation
+
Configuration


Therefore, customers are effectively installing a **business capability**.

20. Multi-Tenant SaaS

This is a commercial enterprise platform.

Multiple customers must use the same platform safely.

Platform



No tenant should ever be able to access another tenant's:

- Agents
- Knowledge
- Credentials
- Executions
- Users
- Logs
- Business data

21. Production Quality

This platform is intended to be sold to businesses.

Therefore, prototypes must eventually evolve into:

High Availability

+

Security

+

Scalability

+

Observability

+

Reliability

+

Auditability

+

Versioning

+

Governance

We should not build a demo that requires a complete rewrite to become production.

The prototype should establish the correct product model and architecture from the beginning.

22. What We Are NOT Building

This clarification is extremely important.

We are NOT building:

A simple chatbot

User → LLM → Response

A simple workflow builder

Node → Node → Node

A prompt management system

Prompt → LLM

An API wrapper

UI → OpenAI

A code generator

Prompt → Python → Execute

We are building:

Enterprise Agent Platform

that combines:

AI

+

Agents

+

Workflows

+

Tools

+

Knowledge

+

Memory

+

Human approval

+

Governance

+

Deployment

+

Observability

23. What Makes This Product Different

The platform should combine the best parts of several categories.

From AI builders

Natural language agent creation.

From workflow platforms

Deterministic workflows and orchestration.

From developer platforms

APIs, SDKs and integrations.

From enterprise automation

Governance, security and approvals.


From AI platforms


Models, RAG, memory and agents.

From marketplaces

Reusable business agents.

The result:





24. The North-Star User Journey

The entire product should ultimately support this journey:

1. I have a business problem.
↓
2. I describe it.
↓
3. The platform designs an agent.
↓
4. I see the workflow.
↓
5. I modify it with AI or visually.
↓
6. I connect my enterprise systems.
↓
7. I give it knowledge.
↓
8. I configure permissions.
↓
9. I test it.
↓

10. I evaluate it.
↓
11. I approve it.
↓
12. I deploy it.
↓
13. It runs in production.
↓
14. I monitor it.
↓
15. I improve it.

If we achieve this experience, we have built the core of the product.

25. The First Prototype

We should NOT attempt to prototype every possible feature.

The first prototype should prove the core experience.

Prototype scenario

Employee Onboarding Agent

User:

Create an employee onboarding agent.

Platform creates:

Employee Created
↓
Get Employee Data
↓
Validate Documents
↓
Search HR Policies

↓
Create Employee Record
↓
Manager Approval
↓
Send Welcome Message

Then demonstrate:

AI Creation
↓
Visual Editing
↓
Tool Configuration
↓
Knowledge Configuration
↓
Human Approval
↓
Test
↓
Live Execution
↓
Execution Trace
↓
Version
↓
Deploy

This is our **golden path**.

26. The First Release

The first product release should allow a customer to:

Create

- Create agent
- Natural-language creation
- Visual creation
- Templates

Configure

- Models
- Tools
- MCP
- Knowledge
- Memory
- Output
- Guardrails
- Human approval

Test

- Manual execution
- Test datasets
- Execution traces

Deploy

- Version
- Staging
- Production
- Rollback

Operate

- Logs
- Monitoring
- Cost
- Latency

- Success/failure

Govern

- RBAC
- Tenant isolation
- Permissions
- Audit

27. Long-Term Vision


The platform eventually evolves from:

Agent Builder

into:

Agent Operating System for Enterprises

Where an organization can have:



Humans and AI agents work together.

Agents can collaborate.

Agents can use enterprise systems.

Agents can escalate to humans.

Agents can operate continuously.

The platform provides the infrastructure, governance and intelligence layer required to make this safe.

28. The Product North Star

The ultimate goal is:

Anyone in an enterprise should be able to describe a business process in natural language and turn it into a secure, governed, observable and production-ready AI agent.

The platform should make AI automation feel as simple as:

Describe



Design



Connect



Test



Deploy

while providing enterprise-grade infrastructure underneath:

Security

Governance

Runtime

Orchestration
Knowledge
Memory
Tools
Evaluation
Observability
Scalability

29. One Sentence For The Team

Every developer should remember this:

We are not building a canvas where users connect AI nodes. We are building a platform where enterprises can turn business intent into trustworthy production AI agents.

That distinction should guide every product and engineering decision.

30. Definition of Success

We will know the product is working when a new customer can come to the platform and, without engineering assistance:

Describe a business problem



Generate an agent



Understand its workflow



Modify it



Connect enterprise tools



Add company knowledge



Test it



Approve it



Deploy it



Watch it operate

That is the product.

Everything else is implementation detail.